Faculté des Sciences et Techniques 2021
Master 2 - ACSYON

TP N°3 - Reinforcement learning

Exercice 1 — TD-Learning and Q-Learning

1. The TD(0) algorithm (described below), allow to estimate the value function for a given policy. An example implementation of this algorithm is given (file `td0_robot.py`) for the recycling robot example. Implement this algorithm in order to evaluate some policies for the parking problem or the gambler problem.

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated
Algorithm parameter: step size $\alpha \in (0, 1]$
Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(terminal) = 0$

Loop for each episode:
  Initialize S
  Loop for each step of episode:
    $A \leftarrow$ action given by π for S
    Take action A, observe R, S'
    $V(S) \leftarrow V(S) + \alpha\big[R + \gamma V(S') - V(S)\big]$
    $S \leftarrow S'$
  until S is terminal

2. Implement the Q-learning algorithm (see the description below) in order to find the optimal policy for the Parking problem or the Gambler problem. An example implementation of this algorithm is given (file `qlearning_robot.py`) for the recycling robot.

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$
Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(terminal, \cdot) = 0$

Loop for each episode:
  Initialize S
  Loop for each step of episode:
    Choose A from S using policy derived from Q (e.g., ε-greedy)
    Take action A, observe R, S'
    $Q(S, A) \leftarrow Q(S, A) + \alpha\big[R + \gamma \max_a Q(S', a) - Q(S, A)\big]$
    $S \leftarrow S'$
  until S is terminal